


Document Data Extraction.

Messy documents in, clean structured data out. This is the core capability behind everything else I build, turning invoices, receipts, certificates and forms into data your systems can actually use.



Raw file




Supplier	✓
Invoice no.	✓
Date	✓
Total	504.00
STRUCTURED DATA	

THE PROBLEM

Someone is retyping every document by hand.

Most businesses still move data from documents into their systems manually. An invoice comes in as a PDF, and someone reads it and types the supplier, the amount, the date and the line items into their accounting tool. A certificate arrives as a scan, and someone copies the reference number and the dates into a spreadsheet. A form gets emailed over, and someone rekeys it into the CRM.



It works, until it does not. Reading and retyping is slow, and everyone makes mistakes eventually. A wrong figure on an invoice or a wrong date on a certificate causes real problems downstream. And every document looks a little different, so there is no clean shortcut. A hundred suppliers means a hundred slightly different layouts. This is the part of a business that quietly eats hours and nobody enjoys.

WHAT I BUILT

It reads documents, so your team does not retype them.

You point the pipeline at whatever is coming in, PDFs, scans, phone photos, or emailed attachments. It reads each one, finds the fields that matter, checks them, and drops clean structured data wherever you need it: a spreadsheet, a CRM, a database, or your accounting software.

THE CORE IDEA

It reads the document, not a template.

A new supplier or a new layout does not break it, because it understands the content on the page rather than matching fixed positions. That is the difference between a system you babysit and one that just works.

AT A GLANCE



Ingests anything

PDFs, scans, photos and email attachments all go in the same door.



Reads any layout

A supplier it has never seen is handled with no template to build.



Checks the data

Validation catches errors before anything is trusted or passed on.



Delivers clean data

Straight into the tools you already use, in the format you already use.

From raw file to clean, checked data.

1

A document comes in

It arrives however it arrives, as a PDF, a scan, a phone photo, or an email attachment.

2

It gets read

Text-based files are read directly. Scans and photos go through OCR or a vision model first, so a crumpled photo of a receipt is handled the same as a clean PDF.

3

The fields get pulled out

The system extracts exactly the things you care about into the same structure every time, so the output slots straight into your systems.

Supplier

Invoice no.

Dates

Line items

Tax

Total

4

It gets checked

✓ CONFIDENT

Line items add up, dates are valid, nothing is missing. It flows straight through.

! UNSURE

Anything it cannot confirm is flagged and sent to a person for a quick check.

5

Clean data goes where it needs to

The finished record is pushed into your spreadsheet, CRM, database, or accounting tool, ready to use.

A WORKED EXAMPLE

One messy invoice, every field in its place.

A supplier invoice arrives as a slightly wonky scanned PDF. To older software it is just a picture with text scattered around. Here is what comes out the other side, the same fields every time, whatever the original looked like.

GOES IN



Scanned PDF

✓ COMES OUT

Supplier	Northgate Electrical Supplies
Invoice number	INV-20481
Invoice date	2026-01-14
Due date	2026-02-13
Line items	3
Subtotal	420.00
VAT (20%)	84.00
Total	504.00

Line items come out as structured rows too, description, quantity and price, not a blob of text.

The same pipeline handles receipts, purchase orders, compliance certificates and forms. In a trades business, that includes the certificates and reports that pile up daily. The document type changes; the approach does not.

WHAT MAKES IT RELIABLE

Accurate, and honest about what it is unsure of.



Reads any layout
A supplier it has never seen is handled with no setup, because there are no templates.



Checks its own work
Validation rules catch the obvious errors, like line items that do not add up.



Knows when it is unsure
Low-confidence results are flagged for a person instead of quietly guessing.



A human stays in the loop
You decide how much gets reviewed and how much flows straight through.



Handles tables, not just fields
Line items, quantities and prices come out as structured rows.



Fits your systems
Output lands in the tool you already use, in the format you already use.

UNDER THE HOOD

There is real engineering behind it. Documents are read with a mix of OCR and vision models for anything scanned or photographed. Extraction runs on Anthropic's Claude, with a fixed output schema so every record comes back in the same shape. Validation rules run over the results before anything is trusted, low-confidence records route to a review queue, and finished data is pushed out through an API into whatever system it needs to reach.



Ingest
Any file



OCR
Vision



Extraction
Claude



Validation
Rules



Delivery
Your systems

WHY THIS BEATS THE OLD WAY

Reads for meaning, not fixed positions.

TEMPLATE OCR

Fixed field positions on the page
Breaks when a layout changes
A new template per new supplier
A slightly rotated scan can fail



THIS PIPELINE

Reads for meaning, not position
Handles layouts it has never seen
No per-supplier setup
Copes with wonky scans and photos

THE RESULTS

PER DOCUMENT

**Minutes →
Seconds**

Read and structured hands-off.

REVIEWED BY A PERSON

**Only the
unsure**

The rest flows straight through.

HANDLED DAILY

Hundreds

Without adding staff to keep up.

Figures shown are representative of the system's design and typical impact. Replace with your own measured results before publishing.

What I can build for you

If your team is retyping data off documents, invoices, receipts, purchase orders, certificates, forms, or anything else that arrives as a file, I can automate it. The pipeline reads whatever comes in, pulls out exactly the fields you need, checks them, and delivers clean data into the systems you already use, with a person reviewing only what actually needs it.